

Lab5 (b): CSC_52002_EP Human motion generation

Advanced deep learning

February 2, 2026




Figure 1: Examples of text-to-motion generation results [7].

In this lab, we explore the field of **text-to-human motion generation**. The goal is to generate a human motion sequence that corresponds to a given text caption, as illustrated in Figure 1.

The lab is structured as follows: first, we examine the **dataset and human motion representation (Section 1)**; next, we study the **diffusion framework (Section 3)**; then, we explore **different architectures (Section 2)**; and finally, we **analyze the results (Section 4)**.

Code 1: Upload the notebook in colab, set a GPU runtime and run cells of the Setup section to setup.

1 Human motion dataset and representation

1.1 HumanML3D dataset

HumanML3D dataset [1] is a 3D human motion-language dataset built from motion capture (mocap) acquisitions. It encompasses a wide range of human actions, including daily activities (e.g., "walking," "jumping"), sports (e.g., "swimming," "playing golf"), acrobatics (e.g., "cartwheel"), and artistic movements (e.g., "dancing").

Question 1: How many motions and descriptions are included in the HumanML3D dataset, and what are their average lengths?

Answer:

1.2 SMPL representation

The SMPL [5] representation is a realistic 3D mesh model of the human body, learned from thousands of 3D body scans. It is based on a set of trainable parameters that capture body shape and pose variations.

Given a set of input parameters, the SMPL model infers a detailed 3D mesh consisting of vertices (points in 3D space that define the shape) and faces (triangular surfaces connecting the vertices to form the body structure).

Question 2: What are the different input parameters of the SMPL model used to infer the vertices, and what do they represent? **Hint:** Check the official code of SMPL models.

Answer:

We cannot directly learn to generate the SMPL parameters directly, we need to use more compact features. We represent a human motion sequence $\mathbf{x} \in \mathbb{R}^{F \times d}$ where F is the number of frames and d the feature dimension. We have $\mathbf{x} = [r_z, \dot{r}_x, \dot{r}_y, \dot{\alpha}, \theta, \mathbf{j}]$, where r_z is the Z (up) coordinate of the pelvis, \dot{r}_x , and \dot{r}_y are the linear velocities of the root, $\dot{\alpha}$ is the angular velocity of the Z angle of the body, θ are the SMPL pose parameters, and \mathbf{j} are the joints positions (computed with the SMPL layer).

Code 2: Complete the missing part in `visualize_smpl.py` to print the dimensions F and d . Then, run the following command to visualize the output video in `smpl.mp4`:

```
python src/visualize_smpl.py
```

2 Model architectures

In this section, we implement three distinct architectures, as shown in Figure 2, inspired by the DiT paper [6]. The paper presents three different methods for conditioning a diffusion model: incontext, AdALN and cross attention. In the code, we implement these architectures using the base class `BaseDiT` located in `src/models/modules/base.py`. The main inference method is `forward`. In the following section, we will implement parts of `cond_mapping`, `backbone` and `output_projection` of each architecture.

2.1 Config A: Incontext

Incontext conditioning is the simplest approach for incorporating conditioning into a transformer-based diffusion model. It consists in appending the conditioning tokens to the main token stream and letting the self-attention blocks mixing them.

Code 3: As shown in Figure 2.A, fill in the missing parts in `src/models/modules/incontext.py`, then run the following command:

```
python src/models/modules/incontext.py
```




Figure 2: Human generation DiT-like architectures.

2.2 Config B: AdaLN

Adaptive Layer Normalization (adaLN) is a more sophisticated conditioning method. It replaces the standard layer normalization layers in transformer blocks with adaptive layer normalization (adaLN). The goal of adaLN is to regress a shift and a scale parameter from the conditioning tokens t and c .

Question 3: Given γ and β the shift and scale parameters, write the AdaLN operation:

$$\text{adaLN}(\mathbf{x}, \gamma, \beta) = \dots \quad (1)$$

Answer:

Code 4: As shown in Figure 2.B, fill in the missing parts in `src/models/modules/adaln.py`, then run the following command:

```
python src/models/modules/adaln.py
```

2.3 Config C: Cross attention

Question 4: Given \mathbf{W}_Q , \mathbf{W}_K , \mathbf{W}_V as the query, key, and value weights, respectively, and d the token dimension, write the cross-attention operation between the input \mathbf{x} and the condition \mathbf{c} :

$$\text{CA}(\mathbf{x}, \mathbf{c}) = \dots \quad (2)$$

Answer:

Code 5: As shown in Figure 2.C, fill in the missing parts in `src/models/modules/cross_attention.py`, then run the following command:

```
python src/models/modules/cross_attention.py
```

Question 5: What is the advantage of using AdaLN for conditioning, particularly in comparison to the cross-attention approach?

Answer:

3 Diffusion framework

In this section, we implement the diffusion loss and the sampling process. **Hint:** You can refer to this diffusion codebook (from page 15) to easily retrieve the relevant equations.

3.1 DDPM

Code 6: DDPM loss. Complete the missing part in `src/training/losses/ddpm.py` to compute the DDPM [4] loss. **Hint:** check Equation 14 of the original DDPM paper [4].

```
python src/training/losses/ddpm.py
```

Code 7: DDPM sampling. Complete the missing part in `src/training/sampler/ddpm.py` to perform the DDPM sampling. **Hint:** check Equation 84 of the diffusion codebook.

Run the following command to generate human motion using DDPM sampling, and visualize the output video in `generation_ddpm_incontext.mp4`: **Hint:** change the value of the random seed to get a different result.

```
python src/generate.py batch_size=1 seed=2 \
diffuser/sampler@diffuser.test_sampler=ddpm
```

3.2 DDIM

Question 6: What are the advantages of DDIM over DDPM sampling, and what is the key difference between them?

Answer:

Code 8: DDIM sampling. Complete the missing part in `src/training/sampler/ddim.py` to perform the DDIM sampling. **Hint:** check Equation 12 of the original DDIM paper [9].

Run the following command to generate human motion using DDIM sampling, and visualize the output video in `generation_ddim_incontext.mp4`: **Hint:** change the value of the random seed to get a different result.

```
python src/generate.py batch_size=1 seed=2 \
diffuser/sampler@diffuser.test_sampler=ddim
```

Question 7: Are there any qualitative differences between the generated samples from the sampling methods? If yes, what are they?

Answer:

4 Result analysis

4.1 Qualitative analysis

Code 9: Run the following commands to generate a sample using each architecture:

```
python src/generate.py batch_size=1 diffuser/network=incontext \
checkpoint_path=./humanml3d-data/checkpoints/incontext.ckpt
```

```
python src/generate.py batch_size=1 diffuser/network=adaln \
checkpoint_path=./humanml3d-data/checkpoints/adaln.ckpt
```

```
python src/generate.py batch_size=1 diffuser/network=cross_attention \
checkpoint_path=./humanml3d-data/checkpoints/cross_attention.ckpt
```

Question 8: Are there any qualitative differences between the generated samples from the different architectures? If yes, what are they?

Answer:

4.2 Quantitative analysis

To quantitatively evaluate the generated motion, we adopt the same approach used in image generation: computing metrics between reference and generated features from an external model. Specifically, we use the TMR [8] feature encoder to calculate the Fréchet Distance (FD_{TMR}) which assesses overall generation quality (similar to FID [3]), and the TMR-Score, which evaluates text-motion coherence (similar to CLIP-Score [2]).

Question 9: What is the main assumption about the reference and generated feature spaces when computing the Fréchet Distance?

Answer:

Code 10: Complete the missing part in `src/metrics/frechet.py` to compute the FD_{TMR} .

Code 11: Complete the missing part in `src/metrics/similarity.py` to compute the TMR-Score.

Bonus 1: Review the code in `src/metrics/motion_text.py` and `src/metrics/prdc.py`. Then, explain the meaning of the following metrics: R1, R2, R3, and PRDC.

Answer:

Code 12: Run the following commands to evaluate each architecture:

```
python src/evaluate.py diffuser/network=incontext \
checkpoint_path=./humanml3d-data/checkpoints/incontext.ckpt
```

```
python src/evaluate.py diffuser/network=adaln \
checkpoint_path=./humanml3d-data/checkpoints/adaln.ckpt
```

```
python src/evaluate.py diffuser/network=cross_attention \
checkpoint_path=./humanml3d-data/checkpoints/cross_attention.ckpt
```

Question 10: Given the computed scores, what can you say?

Answer:

Bonus 2: Here we have evaluated our models on 10×64 samples. How reliable do you think each metric is, particularly FD_{TMR} , and why? What should we do?

Answer:

References

- [1] Chuan Guo, Shihao Zou, Xinxin Zuo, Sen Wang, Wei Ji, Xingyu Li, and Li Cheng. Generating diverse and natural 3d human motions from text. In *CVPR*, 2022.
- [2] Jack Hessel, Ari Holtzman, Maxwell Forbes, Ronan Le Bras, and Yejin Choi. Clipscore: A reference-free evaluation metric for image captioning. In *EMNLP*, 2021.
- [3] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. Gans trained by a two time-scale update rule converge to a local nash equilibrium. *NeurIPS*, 2017.
- [4] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *NeurIPS*, 2020.
- [5] Matthew Loper, Naureen Mahmood, Javier Romero, Gerard Pons-Moll, and Michael J. Black. SMPL: A skinned multi-person linear model. *ACM TOG*, 2015.
- [6] William S Peebles and Saining Xie. Scalable diffusion models with transformers. In *ICCV*, 2022.
- [7] Mathis Petrovich, Michael J Black, and Güл Varol. Temos: Generating diverse human motions from textual descriptions. In *ECCV*, 2022.
- [8] Mathis Petrovich, Michael J Black, and Güл Varol. Tmr: Text-to-motion retrieval using contrastive 3d human motion synthesis. In *ICCV*, 2023.
- [9] Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. *ICLR*, 2021.